

```
import pandas as pd
import numpy as np
import seaborn as sns
sns.set_theme(style='white')
from matplotlib import pyplot as plt
```

Singapore Population Prediction and Enhancement

```
SgData = pd.read_csv('SgData.csv')
SgData
```


	Year	TR	TMR	TFR	SR	MAR
0	2022	4073239	1990212	2083027	955	42.1
1	2021	3986842	1953114	2033728	960	41.8
2	2020	4044210	1977556	2066654	957	41.5
3	2019	4026209	1969382	2056827	957	41.1
4	2018	3994283	1955838	2038445	959	40.8
...
61	1961	1702400	886500	815900	1087	17.9
62	1960	1646400	859600	786800	1093	18.0
63	1959	1587200	830800	756400	1098	18.3
64	1958	1518800	797600	721200	1106	18.5
65	1957	1445929	762760	683169	1117	18.8

[66 rows x 6 columns]

TR = Total Residents TMR = Total Male Residents TFR = Total Female Residents SR = Sex Ratio MAR = Median Age of Residents

a. Year-Population Plot

```
sns.scatterplot(x=SgData['Year'], y=SgData['TR'])
plt.title('Singapore Population 1957-2022')
plt.ylabel('Total Resident')
plt.xlabel('Year')
plt.show()
```



b. Model Fitting

```
sg_train = SgData.loc[SgData['Year']<=2013, :]
sg_test = SgData.loc[SgData['Year']>2013, :]

from sklearn.linear_model import LinearRegression

x1 = sg_train['Year']
y1 = sg_train['TR']
x2 = sg_test['Year']
y2 = sg_test['TR']

x1 = np.array(x1).reshape(-1, 1)
x2 = np.array(x2).reshape(-1, 1)
y1 = np.array(y1).reshape(-1, 1)
y2 = np.array(y2).reshape(-1, 1)

rgs = LinearRegression()
rgs.fit(x1, y1)
y1_pred = rgs.predict(x1)
```

c. i. Best-Fit Line

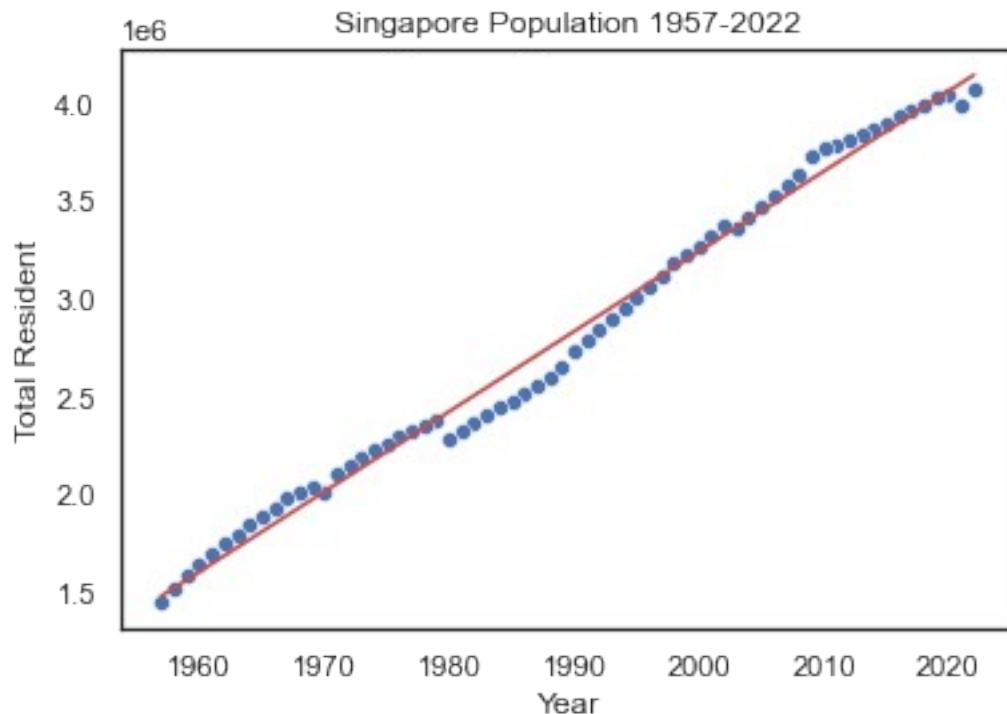
```
print('Slope: ', rgs.coef_)
print('y-Intercept: ', rgs.intercept_)

Slope: [[40975.96590614]]
y-Intercept: [-78707770.39387262]
```

```

sns.scatterplot(x=SgData['Year'], y=SgData['TR'])
temp = SgData['Year'].map(lambda x: rgs.coef_*x+rgs.intercept_)
plt.plot(SgData['Year'], temp, color='r')
plt.title('Singapore Population 1957-2022')
plt.ylabel('Total Resident')
plt.xlabel('Year')
plt.show()

```



c. ii. R^2 Coefficient

```

from sklearn.metrics import mean_squared_error, r2_score

```

```

r2 = r2_score(y1, y1_pred)

```

```

print('R2: ', r2)

```

```

R2: 0.9857883020995728

```

c. iii. Training MSE

```

mse = mean_squared_error(y1, y1_pred)

```

```

print('Train MSE: ', mse)

```

```

Train MSE: 6551722294.557682

```

c. iv. Testing

```

y2_pred = rgs.predict(x2)

```

```

print(np.column_stack((x2, y2_pred)))

```

```

[[2.02200000e+03 4.14563267e+06]
 [2.02100000e+03 4.10465670e+06]
 [2.02000000e+03 4.06368074e+06]
 [2.01900000e+03 4.02270477e+06]
 [2.01800000e+03 3.98172880e+06]
 [2.01700000e+03 3.94075284e+06]
 [2.01600000e+03 3.89977687e+06]
 [2.01500000e+03 3.85880091e+06]
 [2.01400000e+03 3.81782494e+06]]

```

c. v. Testing MSE

```

mse = mean_squared_error(y2, y2_pred)
print('Testing MSE: ', mse)

```

Testing MSE: 2907187388.7105937

d. Prediction

```

X = []
for i in range(2030, 2051):
    X.append(i)
X = np.array(X).reshape(-1, 1)

Y = rgs.predict(X)
rs = np.column_stack((X, Y))
rs = pd.DataFrame(rs)
rs.columns = ['Year', 'Population']
rs

```

	Year	Population
0	2030.0	4.473440e+06
1	2031.0	4.514416e+06
2	2032.0	4.555392e+06
3	2033.0	4.596368e+06
4	2034.0	4.637344e+06
5	2035.0	4.678320e+06
6	2036.0	4.719296e+06
7	2037.0	4.760272e+06
8	2038.0	4.801248e+06
9	2039.0	4.842224e+06
10	2040.0	4.883200e+06
11	2041.0	4.924176e+06
12	2042.0	4.965152e+06
13	2043.0	5.006128e+06
14	2044.0	5.047104e+06
15	2045.0	5.088080e+06
16	2046.0	5.129056e+06
17	2047.0	5.170032e+06
18	2048.0	5.211008e+06
19	2049.0	5.251984e+06
20	2050.0	5.292960e+06

This result is partly reasonable. It is because, in natural population growth principles, there should be competition between samples and the growth rate might naturally decrease to 0. However, in the linear regression model, the growth rate is a constant value. Perhaps it fits the situation in a short term, but it might not precisely reflect the population growth tendency in long term (e.g., >100 years).

e. In the middle of the logistic pattern.

f. Enhancement In this case, we apply a multivariate linear regression. We selected some attributes from the Singapore official statistics data to increase the dimensions of the feature set. Before the formal multivariate linear regression, we conducted a correlation analysis to select the most influencing ones from these additional attributes. To prevent the imbalanced influences of data from different ranges, a normalization phase should be implemented.

```
from sklearn.preprocessing import MinMaxScaler
```

```
sc = MinMaxScaler()
sc = sc.fit(SgData)
tmp = sc.transform(SgData)
tmp = pd.DataFrame(tmp)
tmp.columns = SgData.columns
tmp
```

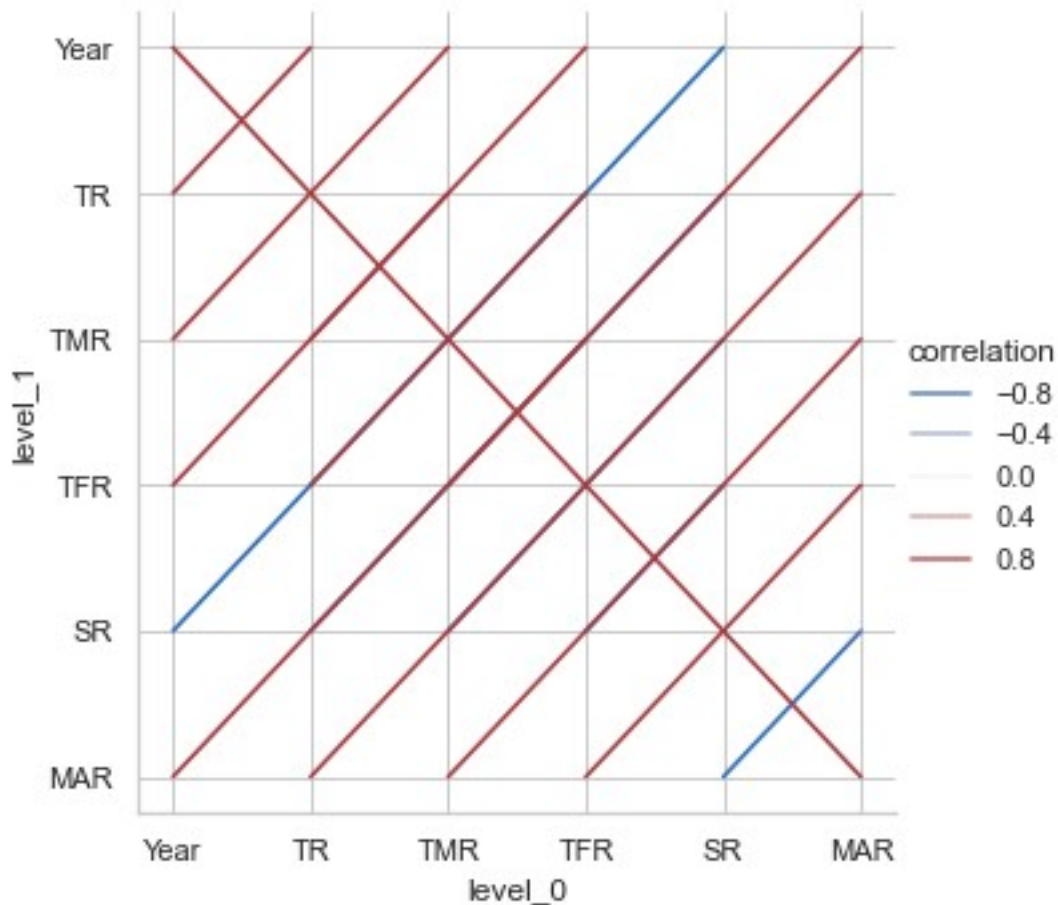
	Year	TR	TMR	TFR	SR	MAR
0	1.000000	1.000000	1.000000	1.000000	0.000000	1.000000
1	0.984615	0.967116	0.969776	0.964783	0.030864	0.987705
2	0.969231	0.988951	0.989689	0.988304	0.012346	0.975410
3	0.953846	0.982100	0.983030	0.981284	0.012346	0.959016
4	0.938462	0.969948	0.971996	0.968152	0.024691	0.946721
...
61	0.061538	0.097617	0.100810	0.094817	0.814815	0.008197
62	0.046154	0.076303	0.078895	0.074030	0.851852	0.012295
63	0.030769	0.053770	0.055432	0.052313	0.882716	0.024590
64	0.015385	0.027736	0.028384	0.027168	0.932099	0.032787
65	0.000000	0.000000	0.000000	0.000000	1.000000	0.045082

[66 rows x 6 columns]

```
corr_mat = tmp.corr().stack().reset_index(name='correlation')
print(corr_mat)
sns.set_theme(style='whitegrid')
sns.relplot(data=corr_mat,
            x='level_0',
            y='level_1',
            palette="vlag",
            hue_norm=(-1, 1),
            hue='correlation',
            kind='line'
            )
```

	level_0	level_1	correlation
0	Year	Year	1.000000
1	Year	TR	0.995049
2	Year	TMR	0.995224
3	Year	TFR	0.994672
4	Year	SR	-0.986310
5	Year	MAR	0.993768
6	TR	Year	0.995049
7	TR	TR	1.000000
8	TR	TMR	0.999865
9	TR	TFR	0.999895
10	TR	SR	-0.985545
11	TR	MAR	0.986491
12	TMR	Year	0.995224
13	TMR	TR	0.999865
14	TMR	TMR	1.000000
15	TMR	TFR	0.999521
16	TMR	SR	-0.985694
17	TMR	MAR	0.986876
18	TFR	Year	0.994672
19	TFR	TR	0.999895
20	TFR	TMR	0.999521
21	TFR	TFR	1.000000
22	TFR	SR	-0.985193
23	TFR	MAR	0.985931
24	SR	Year	-0.986310
25	SR	TR	-0.985545
26	SR	TMR	-0.985694
27	SR	TFR	-0.985193
28	SR	SR	1.000000
29	SR	MAR	-0.963656
30	MAR	Year	0.993768
31	MAR	TR	0.986491
32	MAR	TMR	0.986876
33	MAR	TFR	0.985931
34	MAR	SR	-0.963656
35	MAR	MAR	1.000000

<seaborn.axisgrid.FacetGrid at 0x128d3f730>



TR = Total Residents TMR = Total Male Residents TFR = Total Female Residents SR = Sex Ratio MAR = Median Age of Residents

From the correlation plot above, we can conclude that TMR and SR have the greatest positive influence on TR because the correlation coefficients between these attributes and TR are equal to 1. Therefore, we can add these two attributes to our X matrix to improve the accuracy of the prediction.

```
df = SgData[['Year', 'TR', 'TMR', 'SR']]
```

```
df
```

	Year	TR	TMR	SR
0	2022	4073239	1990212	955
1	2021	3986842	1953114	960
2	2020	4044210	1977556	957
3	2019	4026209	1969382	957
4	2018	3994283	1955838	959
...
61	1961	1702400	886500	1087
62	1960	1646400	859600	1093
63	1959	1587200	830800	1098
64	1958	1518800	797600	1106

```
65  1957  1445929   762760  1117
```

```
[66 rows x 4 columns]
```

```
sc = MinMaxScaler()
sc = sc.fit(df)
train = df.loc[df['Year']<=2013, :]
test = df.loc[df['Year']>2013, :]
tmp_1 = sc.transform(train)
tmp_2 = sc.transform(test)
tmp_1 = pd.DataFrame(tmp_1)
tmp_2 = pd.DataFrame(tmp_2)
tmp_1.columns = df.columns
tmp_2.columns = df.columns

# training
X = tmp_1[['Year', 'TMR', 'SR']]
Y = tmp_1['TR']

rgs = LinearRegression()
rgs.fit(X, Y)
Y_pred = rgs.predict(X)

rs = np.column_stack((np.array(X['Year']),
                        np.array(Y_pred),
                        np.array(X[['TMR', 'SR']])))
rs = sc.inverse_transform(rs)
rs = pd.DataFrame(rs)
rs.columns = df.columns
rs
```

	Year	TR	TMR	SR
0	2013.0	3.824059e+06	1891504.0	968.0
1	2012.0	3.800450e+06	1880046.0	970.0
2	2011.0	3.775927e+06	1868170.0	972.0
3	2010.0	3.761988e+06	1861133.0	974.0
4	2009.0	3.727567e+06	1844732.0	976.0
5	2008.0	3.637403e+06	1802992.0	980.0
6	2007.0	3.578670e+06	1775477.0	982.0
7	2006.0	3.520707e+06	1748242.0	983.0
8	2005.0	3.462875e+06	1721139.0	985.0
9	2004.0	3.407377e+06	1695031.0	986.0
10	2003.0	3.361518e+06	1673401.0	988.0
11	2002.0	3.386488e+06	1684295.0	992.0
12	2001.0	3.331487e+06	1658558.0	995.0
13	2000.0	3.280524e+06	1634667.0	998.0
14	1999.0	3.238530e+06	1614804.0	1000.0
15	1998.0	3.189700e+06	1591816.0	1002.0
16	1997.0	3.133979e+06	1565750.0	1005.0
17	1996.0	3.078989e+06	1540018.0	1008.0
18	1995.0	3.023564e+06	1514015.0	1010.0

19	1994.0	2.970536e+06	1489180.0	1013.0
20	1993.0	2.917084e+06	1464223.0	1017.0
21	1992.0	2.863346e+06	1439063.0	1020.0
22	1991.0	2.810176e+06	1414307.0	1025.0
23	1990.0	2.750348e+06	1386291.0	1027.0
24	1989.0	2.653647e+06	1341275.0	1027.0
25	1988.0	2.601390e+06	1316577.0	1027.0
26	1987.0	2.554182e+06	1294187.0	1027.0
27	1986.0	2.517153e+06	1276522.0	1028.0
28	1985.0	2.479421e+06	1258464.0	1028.0
29	1984.0	2.438695e+06	1239109.0	1029.0
30	1983.0	2.399695e+06	1220471.0	1029.0
31	1982.0	2.357230e+06	1200321.0	1030.0
32	1981.0	2.313951e+06	1179799.0	1031.0
33	1980.0	2.270090e+06	1159011.0	1032.0
34	1979.0	2.395604e+06	1216300.0	1042.0
35	1978.0	2.364622e+06	1201400.0	1043.0
36	1977.0	2.335549e+06	1187300.0	1043.0
37	1976.0	2.302161e+06	1171300.0	1044.0
38	1975.0	2.270524e+06	1156100.0	1045.0
39	1974.0	2.236262e+06	1139700.0	1046.0
40	1973.0	2.198001e+06	1121400.0	1046.0
41	1972.0	2.155207e+06	1101100.0	1047.0
42	1971.0	2.113289e+06	1081200.0	1048.0
43	1970.0	2.004674e+06	1030809.0	1049.0
44	1969.0	2.041169e+06	1046900.0	1052.0
45	1968.0	2.014153e+06	1034100.0	1057.0
46	1967.0	1.978920e+06	1017400.0	1060.0
47	1966.0	1.933283e+06	995800.0	1061.0
48	1965.0	1.886142e+06	973800.0	1066.0
49	1964.0	1.838660e+06	951500.0	1069.0
50	1963.0	1.791239e+06	929300.0	1073.0
51	1962.0	1.748282e+06	909500.0	1082.0
52	1961.0	1.698954e+06	886500.0	1087.0
53	1960.0	1.640937e+06	859600.0	1093.0
54	1959.0	1.578922e+06	830800.0	1098.0
55	1958.0	1.506810e+06	797600.0	1106.0
56	1957.0	1.430639e+06	762760.0	1117.0

```

r2 = r2_score(rs['Year'], train['Year'])
mse = mean_squared_error(rs['Year'], train['Year'])
print('Training MSE: ', mse)
print('Training R2: ', r2)

```

```

Training MSE:  0.0
Training R2:  1.0

```

```

# testing
X = tmp_2[['Year', 'TMR', 'SR']]
Y = tmp_2['TR']

```

```

rgs = LinearRegression()
rgs.fit(X, Y)
Y_pred = rgs.predict(X)

rs = np.column_stack((np.array(X['Year']),
                        np.array(Y_pred),
                        np.array(X[['TMR', 'SR']])))
rs = sc.inverse_transform(rs)
rs = pd.DataFrame(rs)
rs.columns = df.columns
rs

```

	Year	TR	TMR	SR
0	2022.0	4.073281e+06	1990212.0	955.0
1	2021.0	3.986846e+06	1953114.0	960.0
2	2020.0	4.044114e+06	1977556.0	957.0
3	2019.0	4.025991e+06	1969382.0	957.0
4	2018.0	3.994519e+06	1955838.0	959.0
5	2017.0	3.965867e+06	1943545.0	961.0
6	2016.0	3.933324e+06	1929526.0	963.0
7	2015.0	3.903308e+06	1916628.0	965.0
8	2014.0	3.870317e+06	1902410.0	967.0

```

r2 = r2_score(rs['Year'], test['Year'])
mse = mean_squared_error(rs['Year'], test['Year'])
print('Training MSE: ', mse)
print('Training R2: ', r2)

```

```

Training MSE:  5.744309809396025e-27
Training R2:  1.0

```

The testing result of multivariate regression has a 0-like MSE and its R^2 coefficient equals to 1, which is much more precise than the previous simple linear regression with [Testing MSE, R^2] = [2907187388.7105937, 0.9857883020995728]. Therefore, multivariate linear regression can be considered an effective enhancement method.

Chinese Population Prediction

a. Load Dataset

```

CNData = pd.read_csv('CNData.csv')
CNData

```

	Year	PPL
0	1960	667070000
1	1961	660330000
2	1962	665770000
3	1963	682335000
4	1964	698355000
...
57	2017	1396215000
58	2018	1402760000


```
59 2019 1407745000
60 2020 1411100000
61 2021 1412360000
```

[62 rows x 2 columns]

PPL = Population

b. Year-Population Plot

```
sns.set_theme(style='white')
sns.scatterplot(x=CNData['Year'], y=CNData['PPL'])
plt.title('Chinese Population 1960-2021')
plt.ylabel('Total Population')
plt.xlabel('Year')
plt.show()
```



c. Model Training

```
cn_train = CNData.loc[CNData['Year']<=2013, :]  
cn_test = CNData.loc[CNData['Year']>2013, :]  
  
x1 = np.array(cn_train['Year']).reshape(-1, 1)  
x2 = np.array(cn_test['Year']).reshape(-1, 1)  
y1 = np.array(cn_train['PPL']).reshape(-1, 1)  
y2 = np.array(cn_test['PPL']).reshape(-1, 1)
```


```
rgs = LinearRegression()
rgs.fit(x1, y1)
y1_pred = rgs.predict(x1)
```

d. i. Best-Fit Line

```
print('Slope: ', rgs.coef_)
print('y-Intercept: ', rgs.intercept_)
```

```
Slope: [[14006789.21288355]]
y-Intercept: [-2.67690896e+10]
```

```
temp = CNDData['Year'].map(lambda x: rgs.coef_*x+rgs.intercept_)
sns.scatterplot(x=CNDData['Year'], y=CNDData['PPL'])
plt.plot(CNDData['Year'], temp, color='r')
plt.title('Chinese Population 1960-2021')
plt.ylabel('Total Population')
plt.xlabel('Year')
plt.show()
```



d. ii. R^2 Coefficient

```
r2 = r2_score(y1, y1_pred)
print('R2: ', r2)
```

```
R2: 0.9832476129933141
```

d. iii. Training MSE

```
mse = mean_squared_error(y1, y1_pred)
print('Training MSE: ', mse)
```

Training MSE: 811985540736408.5

d. iv. Testing

```
y2_pred = rgs.predict(x2)
```

d. v. Testing MSE

```
mse = mean_squared_error(y2, y2_pred)
print('Testing MSE: ', mse)
```

Testing MSE: 9065457889543732.0

e. Prediction of 2030-2050

```
X = []
for i in range(2030, 2051):
    X.append(i)
X = np.array(X).reshape(-1, 1)

Y = rgs.predict(X)
rs = np.column_stack((X, Y))
rs = pd.DataFrame(rs)
rs.columns = ['Year', 'Population']
rs
```

	Year	Population
0	2030.0	1.664692e+09
1	2031.0	1.678699e+09
2	2032.0	1.692706e+09
3	2033.0	1.706713e+09
4	2034.0	1.720720e+09
5	2035.0	1.734726e+09
6	2036.0	1.748733e+09
7	2037.0	1.762740e+09
8	2038.0	1.776747e+09
9	2039.0	1.790754e+09
10	2040.0	1.804760e+09
11	2041.0	1.818767e+09
12	2042.0	1.832774e+09
13	2043.0	1.846781e+09
14	2044.0	1.860788e+09
15	2045.0	1.874794e+09
16	2046.0	1.888801e+09
17	2047.0	1.902808e+09
18	2048.0	1.916815e+09
19	2049.0	1.930821e+09
20	2050.0	1.944828e+09

This prediction is not reasonable. From the figure plotted in the previous stage, it can be easily found that the growth rate of the Chinese population is descending. The predicted growth rate (slope) is a large constant because the dataset used includes the phases of acceleration of Chinese population growth. It does not match the current situation in China.

f. Logistic pattern.

g. As mentioned earlier, multivariate linear regression can be used to improve the precision of the prediction. However, due to the lack of further information, we did not verify this enhancement method by predicting the Chinese population.